

Understanding of AlexNet Architecture

PREPARED BY: SUBASH SAH

Introduction

Introduced by Alex Krizhevsky et al. in 2012.

Won the ImageNet Large Scale Visual Recognition Challenge (ILSVRC 2012).

Marked a major breakthrough in deep learning and computer vision.

Key Achievements

Reduced classification error rate from nearly 26% to 15%.

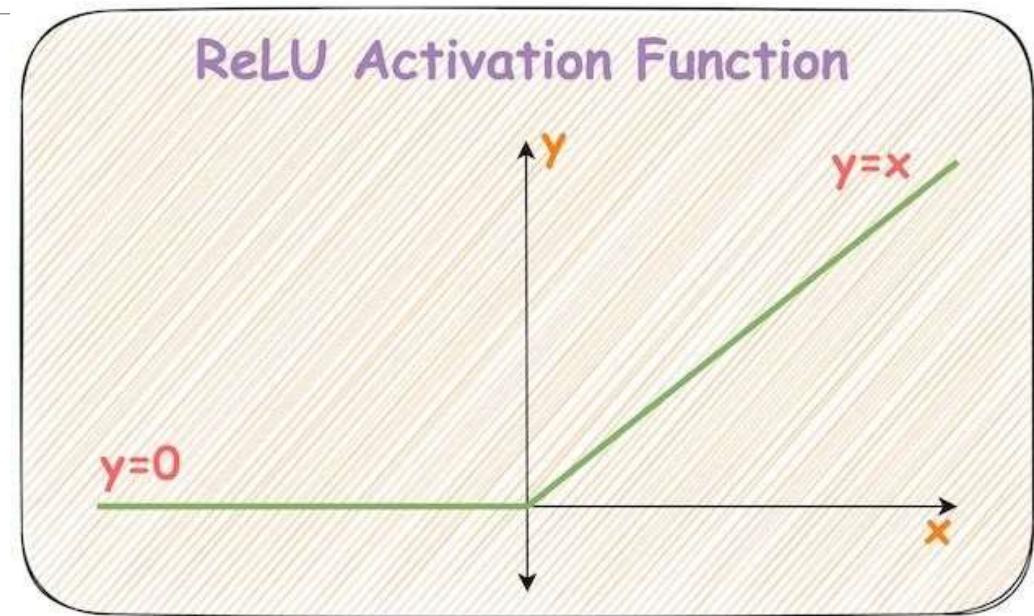
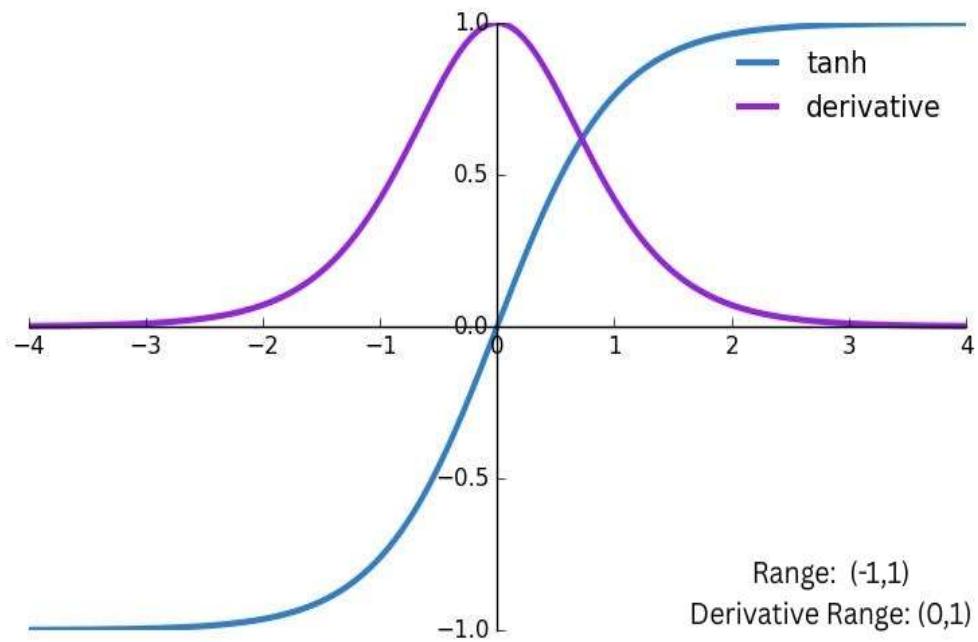
Used **ReLU** activation instead of sigmoid/tanh.

First CNN trained on GPUs – huge performance boost.

Introduced **Dropout** and **Data Augmentation** to prevent overfitting.

Dropout: A regularization technique that randomly deactivates a fraction of neurons during training to prevent overfitting and improve model generalization.

Data Augmentation: A method of artificially increasing the training dataset by applying transformations (like rotation, flipping, cropping or scaling) to existing images to improve model robustness.



Disadvantages of Tanh:

Saturating function -> Creates Vanishing Gradient Problem

Computationally expensive -> Involves exponential calculation

Advantages of ReLU over Tanh:

Non-saturating in +ve region.

Computationally inexpensive.

Converges faster than sigmoid and tanh.

Some terms before moving to architecture:

Convolution Operation

Padding

Strides

Pooling layer

Convolution Operation

Image

| | | | | | |
|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

*

Filter
(Horizontal)

| | | |
|----|----|----|
| -1 | -1 | -1 |
| 0 | 0 | 0 |
| 1 | 1 | 1 |

=

feature map

| | | | |
|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 |
| 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 |
| 0 | 0 | 0 | 0 |

3*3

4*4

6*6
n*n

m*m

(n-m+1)*(n-m+1)

The diagram illustrates a convolution operation. An input image of size 6x6 (6*6, n*n) is multiplied by a horizontal filter of size 3x3 (3*3, m*m). The result is a feature map of size 4x4 (4*4, (n-m+1)*(n-m+1)). The diagram shows the receptive field of each output unit in the feature map, which corresponds to a 3x3 kernel step over the input. The filter values are -1, -1, -1 on top, 0, 0, 0 in the middle, and 1, 1, 1 at the bottom. The resulting feature map has values 0, 0, 0, 0 in the first row and 255, 255, 255, 255 in the second row, with the third and fourth rows being zero.

For RGB Image:

$$\begin{matrix} 6*6*3 \\ n*n*c \end{matrix} * \begin{matrix} 3*3*3 \\ m*m*c \end{matrix} = \begin{matrix} 4*4 \\ (n-m+1)(n-m+1) \end{matrix}$$

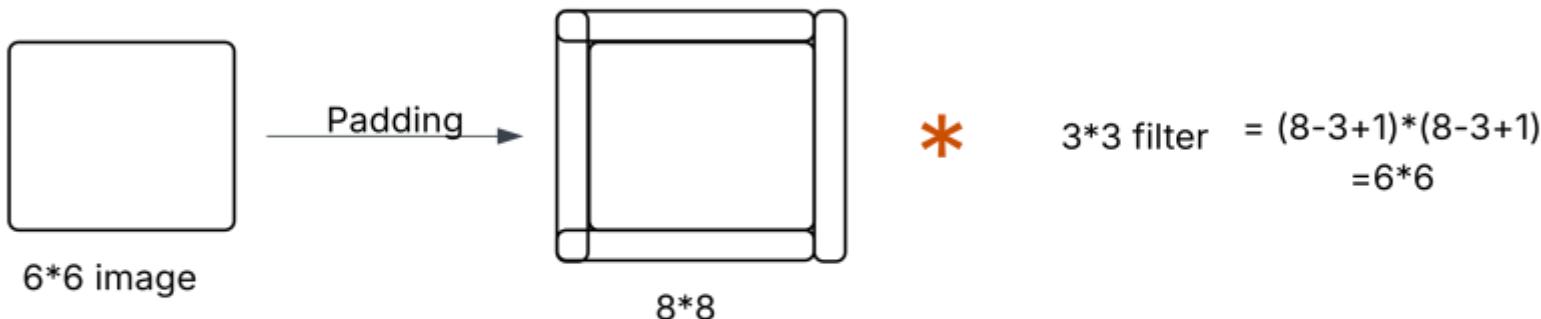
For multiple filters:

$$\begin{matrix} 6*6*3 \\ * \end{matrix} \begin{matrix} 3*3*3 \text{ (horizontal edge)} \\ 3*3*3 \text{ (vertical edge)} \end{matrix} = \begin{matrix} 4*4 \\ 4*4 \end{matrix} \left. \right\} 4*4*2$$

If 10 filters then $4*4*10$ feature map and so on.

Padding

After applying filter, the size of the image is reduced. So, padding is done so that the size doesn't get reduced. In this, before applying filter, we add padding (normally zero padding) to the image and then apply filter so that the size doesn't get reduced.



Dimension of feature map after padding is:
 $n \times n * m \times m = (n+2p-m+1) \times (n+2p-m+1)$

Keras provide two types of padding:

1. Valid: image size is reduced.
2. same: image size remains same

Stride

When not mentioned, the value of stride is 1 by default.

If strides = (2,2), then the sliding window slides 2 units horizontal as well as 2 units vertical.

After applying stride (along with padding), the dimension of feature map will be:

$$n*n \text{ * } m*m = \left(\frac{n+2p-m}{s} + 1\right) * \left(\frac{n+2p-m}{s} + 1\right)$$

Pooling layer

Pooling is a way to down sample the feature map.

Types of pooling:

- Max pooling
- Min pooling
- Avg pooling
- Global pooling

| | | | |
|---|---|---|---|
| 3 | 1 | 1 | 3 |
| 2 | 5 | 0 | 2 |
| 1 | 4 | 2 | 1 |
| 4 | 7 | 2 | 4 |

Pooling →

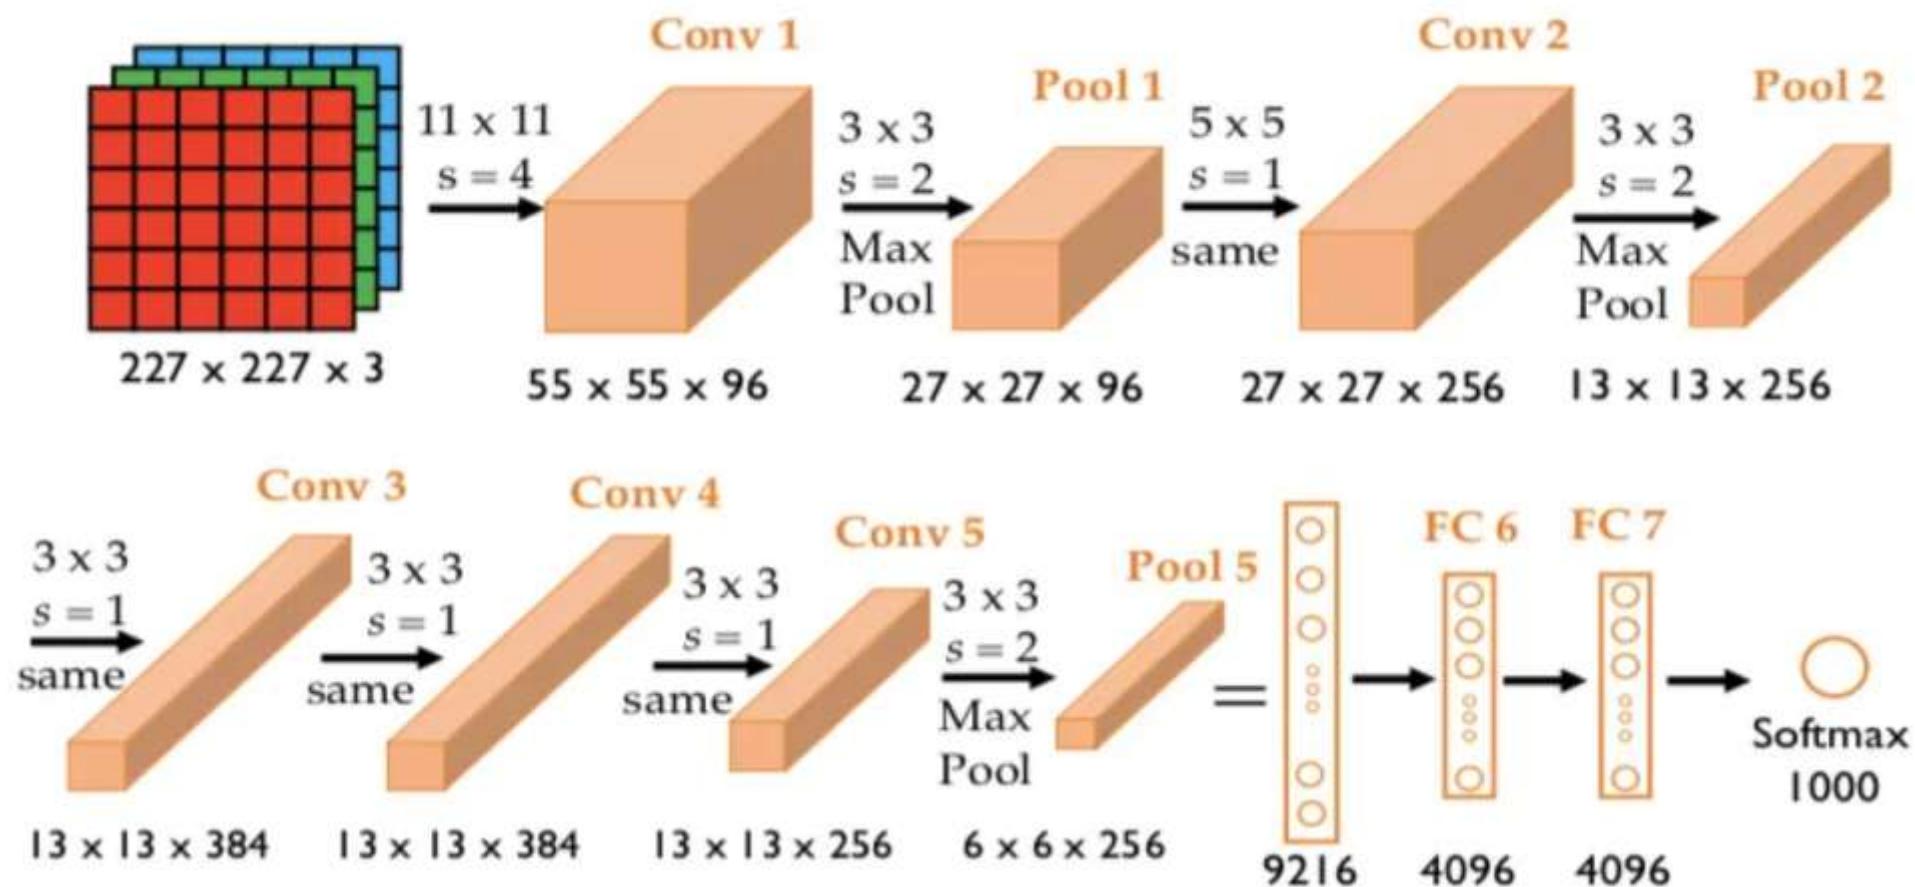
size = (2,2)

stride = (2,2)

type = max

$$\begin{bmatrix} 5 & 3 \\ 7 & 4 \end{bmatrix}$$

Architecture



```
model = Sequential()

# Convolution layer 1
model.add(Conv2D(filters=96,kernel_size=(11,11),strides=(4,4),padding='valid',activation='relu',input_shape=(227,227,3)))
# Pooling layer
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))
# Convolution layer 2
model.add(Conv2D(filters=256,kernel_size=(5,5),strides=(1,1),padding='same',activation='relu'))
# Pooling layer 2
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))
# Convolution layer 3
model.add(Conv2D(filters=384,kernel_size=(3,3),strides=(1,1),padding='same',activation='relu'))
# Convolution layer 4
model.add(Conv2D(filters=384,kernel_size=(3,3),strides=(1,1),padding='same',activation='relu'))
# Convolution layer 5
model.add(Conv2D(filters=256,kernel_size=(3,3),strides=(1,1),padding='same',activation='relu'))
# Pooling layer 3
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))

# Flatten
model.add(Flatten())

# Dense layer 1
model.add(Dense(4096,activation='relu'))
# Dense layer 2
model.add(Dense(4096,activation='relu'))
# Output layer
model.add(Dense(1000,activation='softmax'))

model.summary()
```

| Layer (type) | Output shape | Param # |
|--------------------------------|---------------------|------------|
| conv2d (Conv2D) | (None, 55, 55, 96) | 34,944 |
| max_pooling2d (MaxPooling2D) | (None, 27, 27, 96) | 0 |
| conv2d_1 (Conv2D) | (None, 27, 27, 256) | 614,656 |
| max_pooling2d_1 (MaxPooling2D) | (None, 13, 13, 256) | 0 |
| conv2d_2 (Conv2D) | (None, 13, 13, 384) | 885,120 |
| conv2d_3 (Conv2D) | (None, 13, 13, 384) | 1,327,488 |
| conv2d_4 (Conv2D) | (None, 13, 13, 256) | 884,992 |
| max_pooling2d_2 (MaxPooling2D) | (None, 6, 6, 256) | 0 |
| flatten (Flatten) | (None, 9216) | 0 |
| dense (Dense) | (None, 4096) | 37,752,832 |
| dense_1 (Dense) | (None, 4096) | 16,781,312 |
| dense_2 (Dense) | (None, 1000) | 4,097,000 |

Total params: 62,378,344 (237.95 MB)

Trainable params: 62,378,344 (237.95 MB)

Non-trainable params: 0 (0.00 B)

Explanation for parameters

Layer 1: Conv2D(96 filters, 11×11 kernel, input 3 channels)

$$\gg (11 \times 11 \times 3) \times 96 + 96 = 34,944$$

Layer 2: Conv2D(256 filters, 5×5 kernel, input 96 channels)

$$\gg (5 \times 5 \times 96) \times 256 + 256 = 614,656$$

Layer 3: Conv2D(384 filters, 3×3 kernel, input 256 channels)

$$\gg (3 \times 3 \times 256) \times 384 + 384 = 884,992 + 384 = 885,120$$

Layer 4: Conv2D(384 filters, 3×3 kernel, input 384 channels)

$$\gg (3 \times 3 \times 384) \times 384 + 384 = 1,327,104 + 384 = 1,327,488$$

Layer 5: Conv2D(256 filters, 3×3 kernel, input 384 channels)

$$\gg (3 \times 3 \times 384) \times 256 + 256 = 884,736 + 256 = 884,992$$

Fully connected(Dense) layers

Dense($9216 \rightarrow 4096$):

$$(9216 \times 4096) + 4096 = 37,752,832$$

Dense($4096 \rightarrow 4096$):

$$(4096 \times 4096) + 4096 = 16,781,312$$

Dense($4096 \rightarrow 1000$):

$$(4096 \times 1000) + 1000 = 4,097,000$$

Total: $34,944 + 614,656 + 885,120 + 1,327,488 + 884,992 + 37,752,832 + 16,781,312 + 4,097,000$

$$= 62,378,344$$

Evolution on CNN Architectures on ImageNet

| Year | Model | Error Rate |
|-------|--------------------------------------|-------------|
| 1998 | LeNet-5 | ~25% |
| 2012 | AlexNet | 15.3% |
| 2013 | ZFNet | 11.7% |
| 2014 | VGGNet | 7.3% |
| 2014 | GoogLeNet (Inception v1) | 6.7% |
| 2015 | ResNet | 3.6% |
| 2016 | Inception-v3 / Inception-v4 | 3.5% – 3.1% |
| 2017 | DenseNet | 3.5% |
| 2019 | EfficientNet | <2.0% |
| 2021+ | Vision Transformers (ViT, Swin, etc) | <1.0% |

Thank You