

Attention Is All You Need

By: Subash Sah

Why Transformers?

- RNNs/LSTMs are slow and sequential
- Hard to capture long-range dependencies
- Transformers allow full parallel processing
- Much better performance in translation and NLP

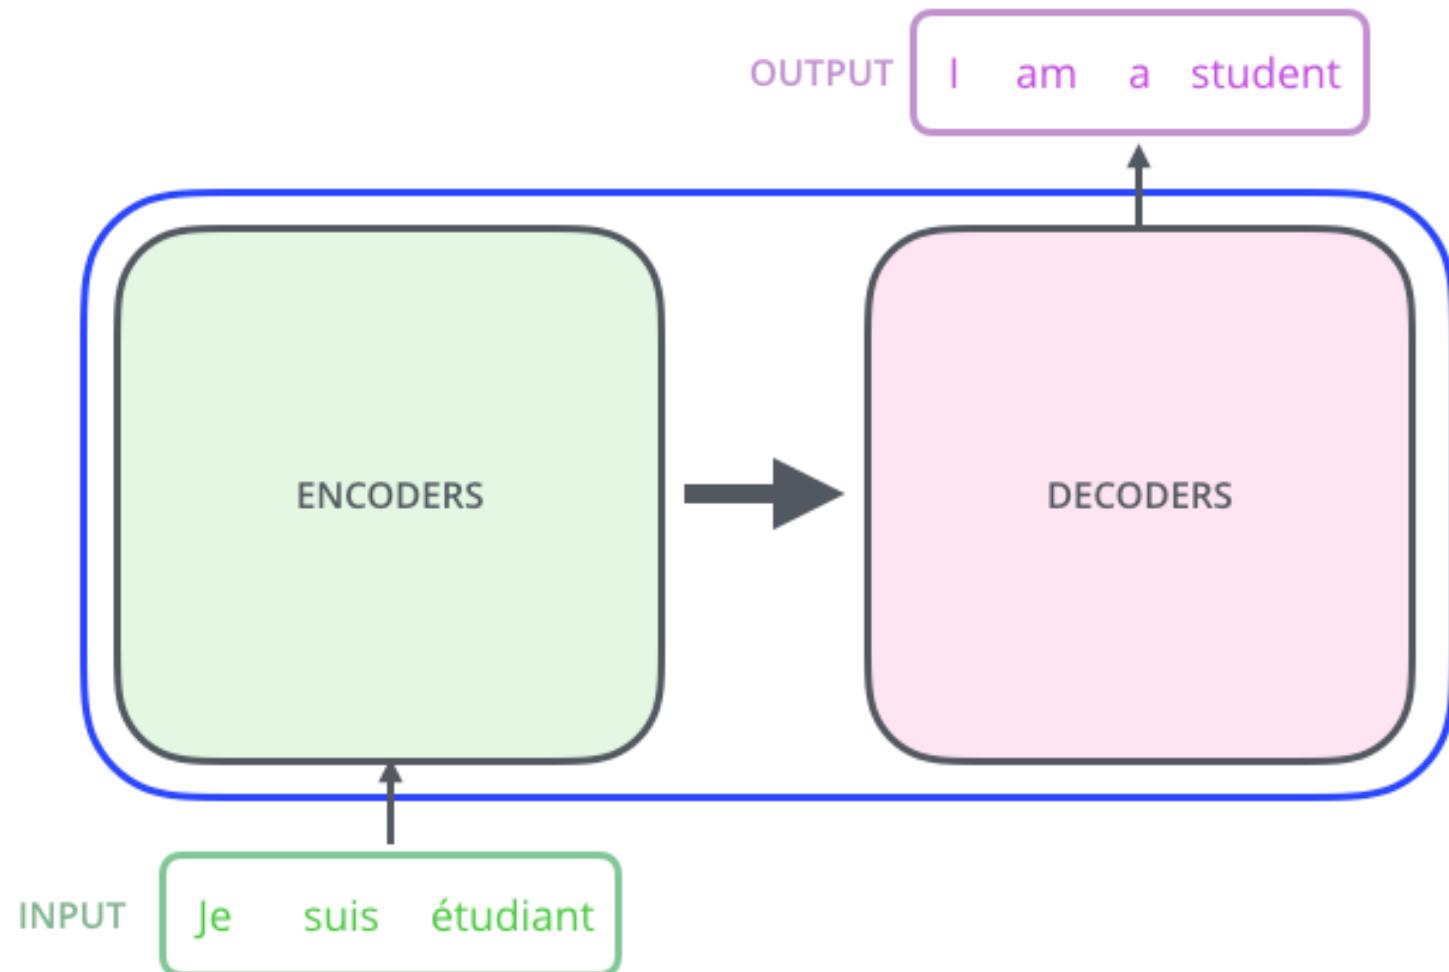
What Makes Transformers Different

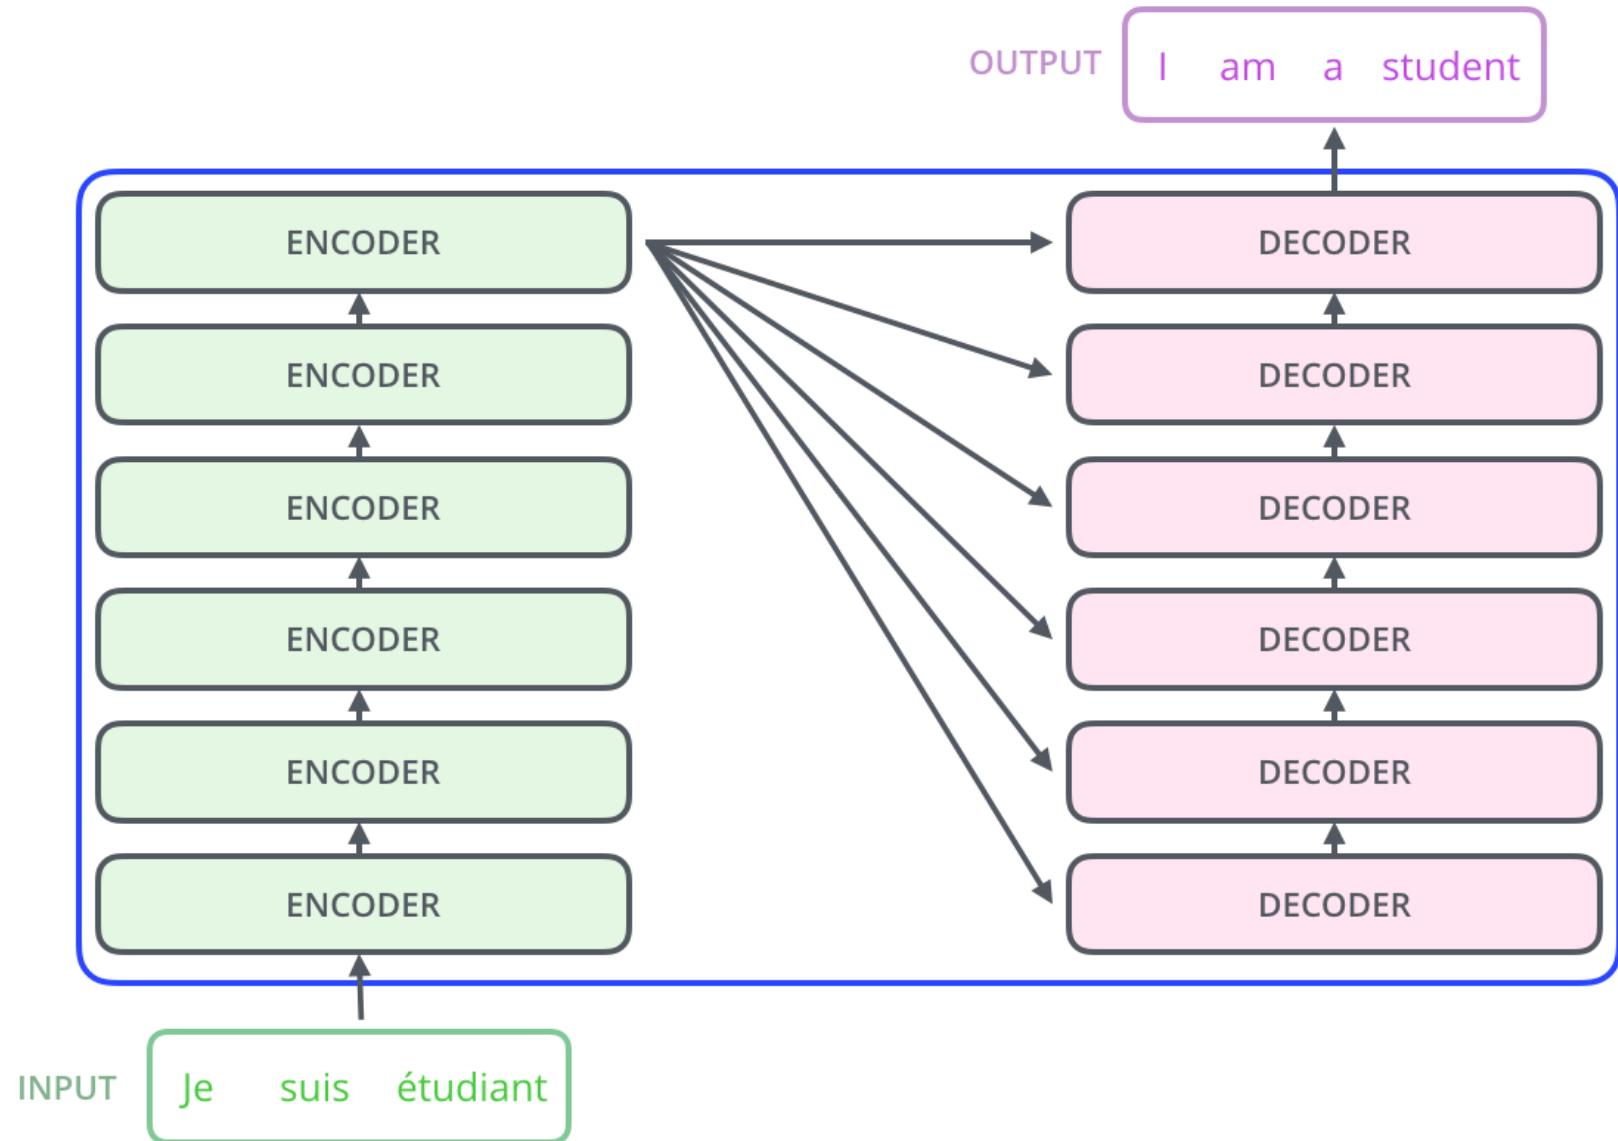
- No recurrence
- Use self-attention instead
- Handle whole sequences at once
- Scale well with data and model size
- Foundation for modern models like BERT and GPT

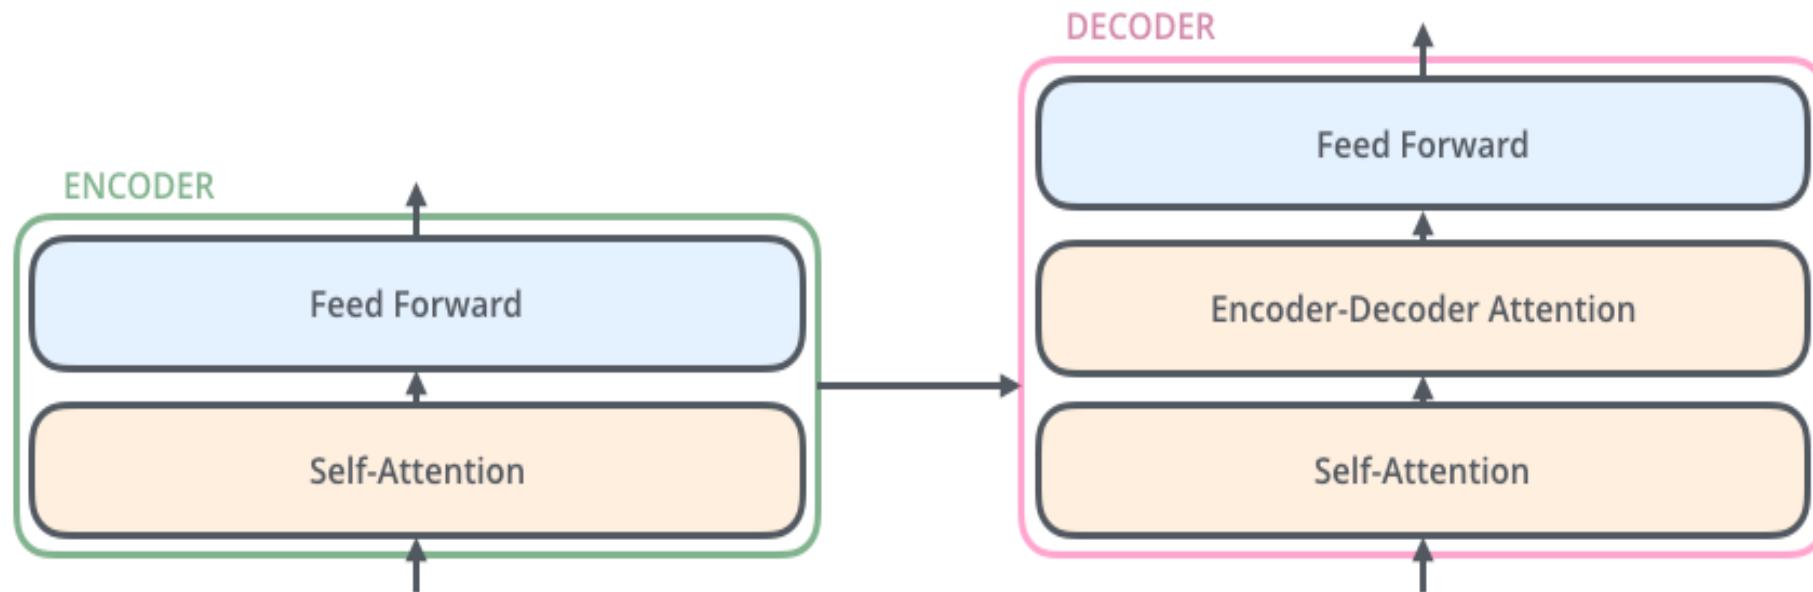
Architecture of Transformer

A High-Level Look

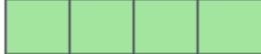




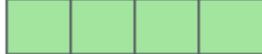




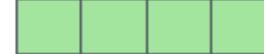
Bringing Tensors/Vectors into the Picture

x_1 

Je

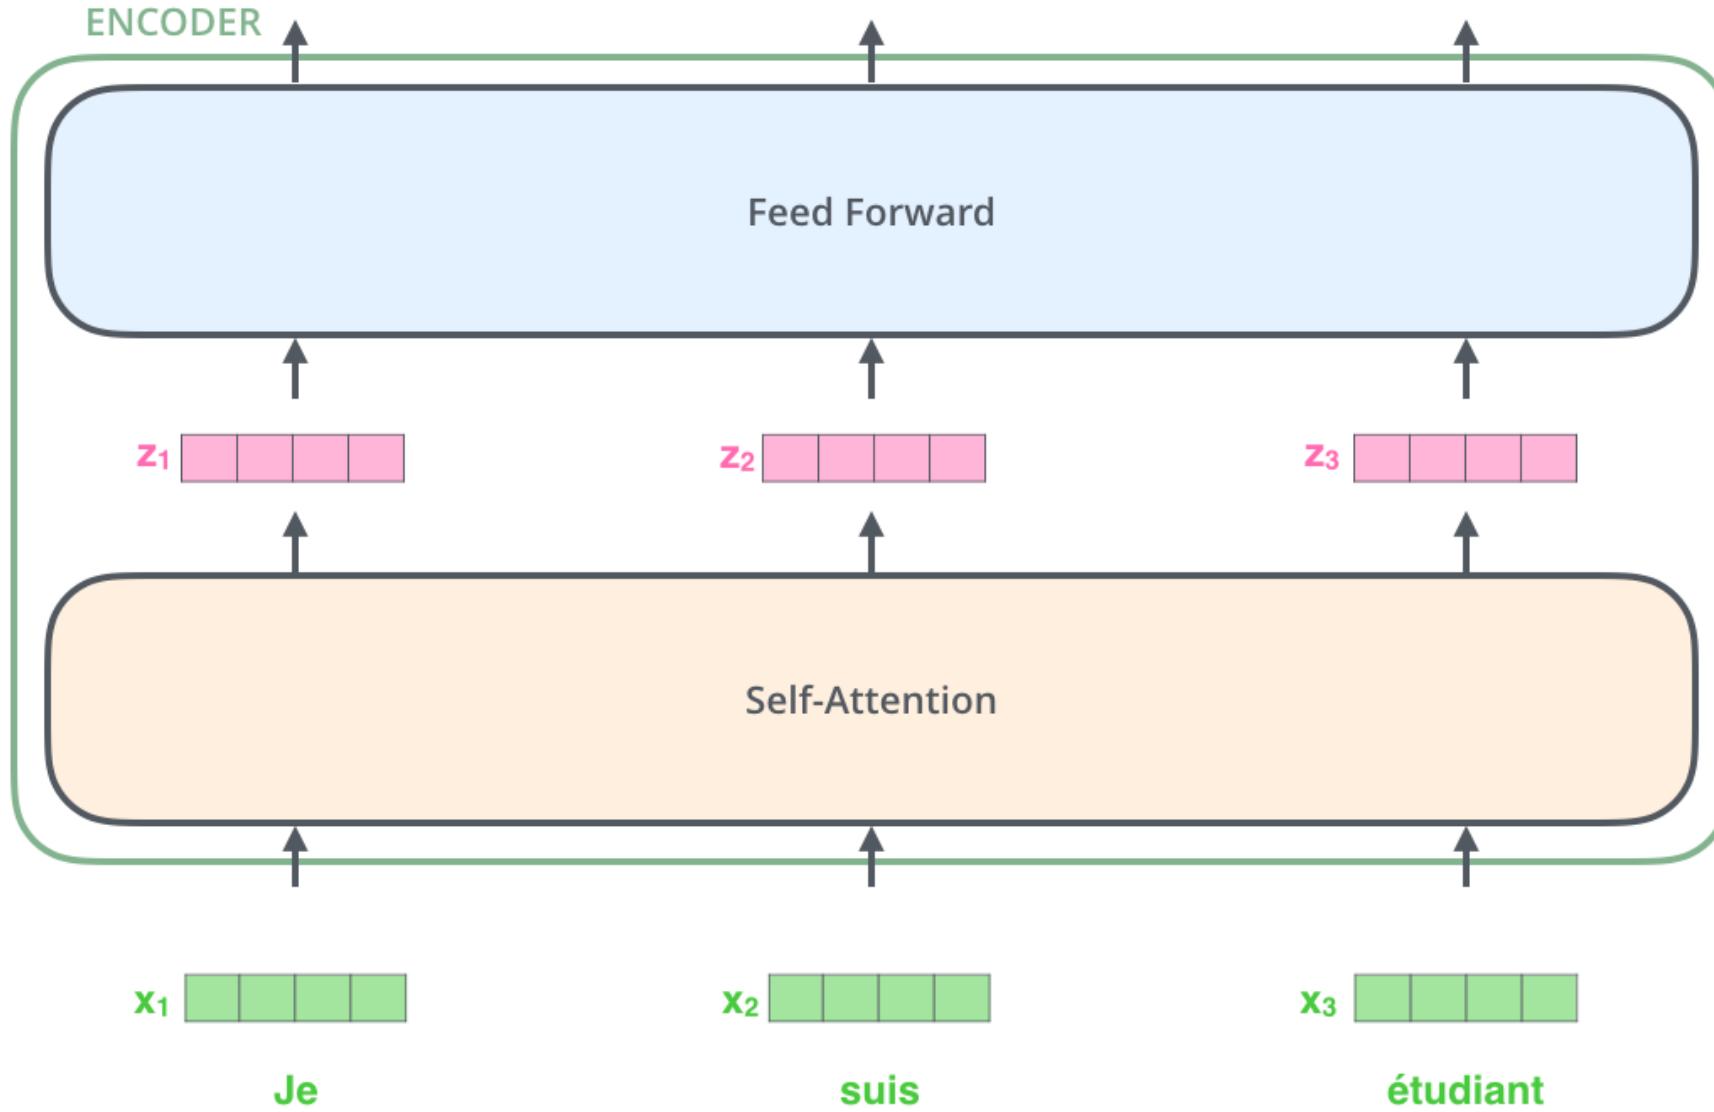
x_2 

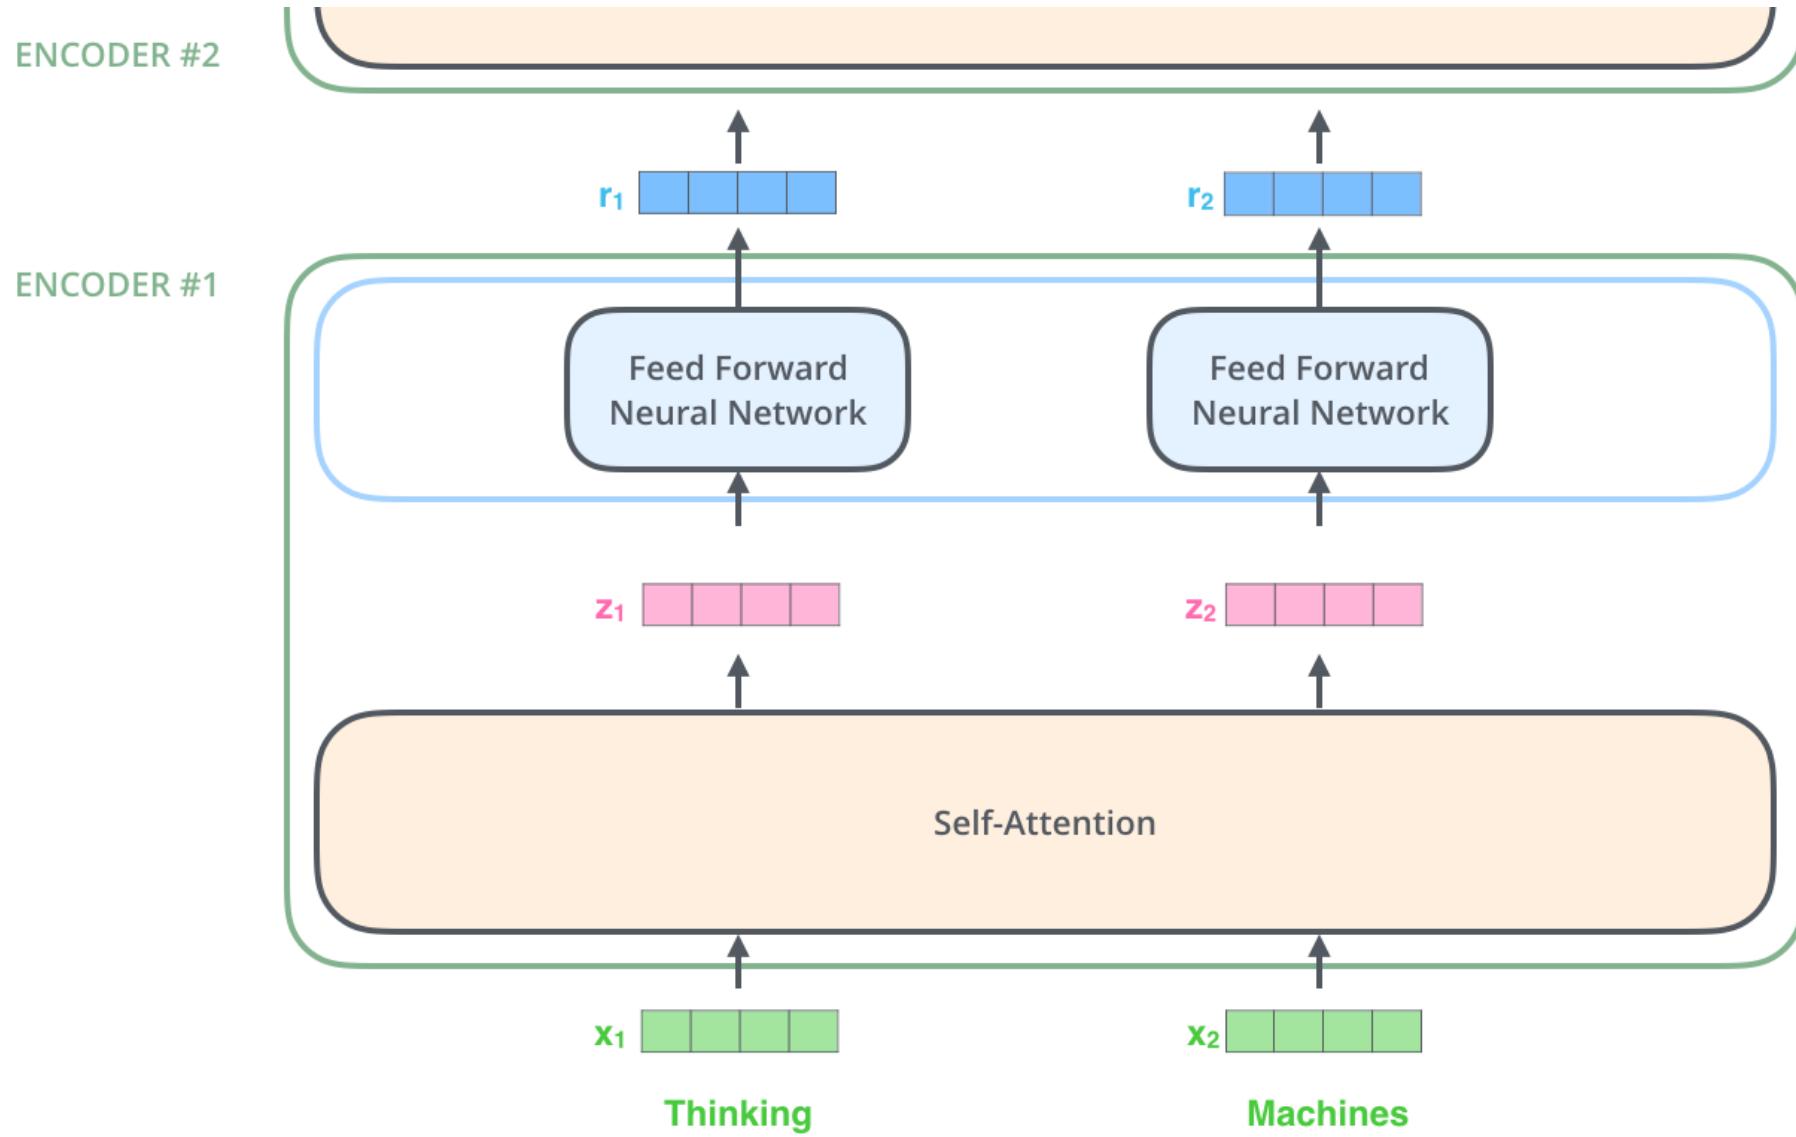
suis

x_3 

étudiant

Vector 512 dimensions





Self-Attention at High Level

- Self-attention is a mechanism where each word in a sentence looks at every other word (including itself) to decide which ones are important for understanding its meaning.

Why we need it?

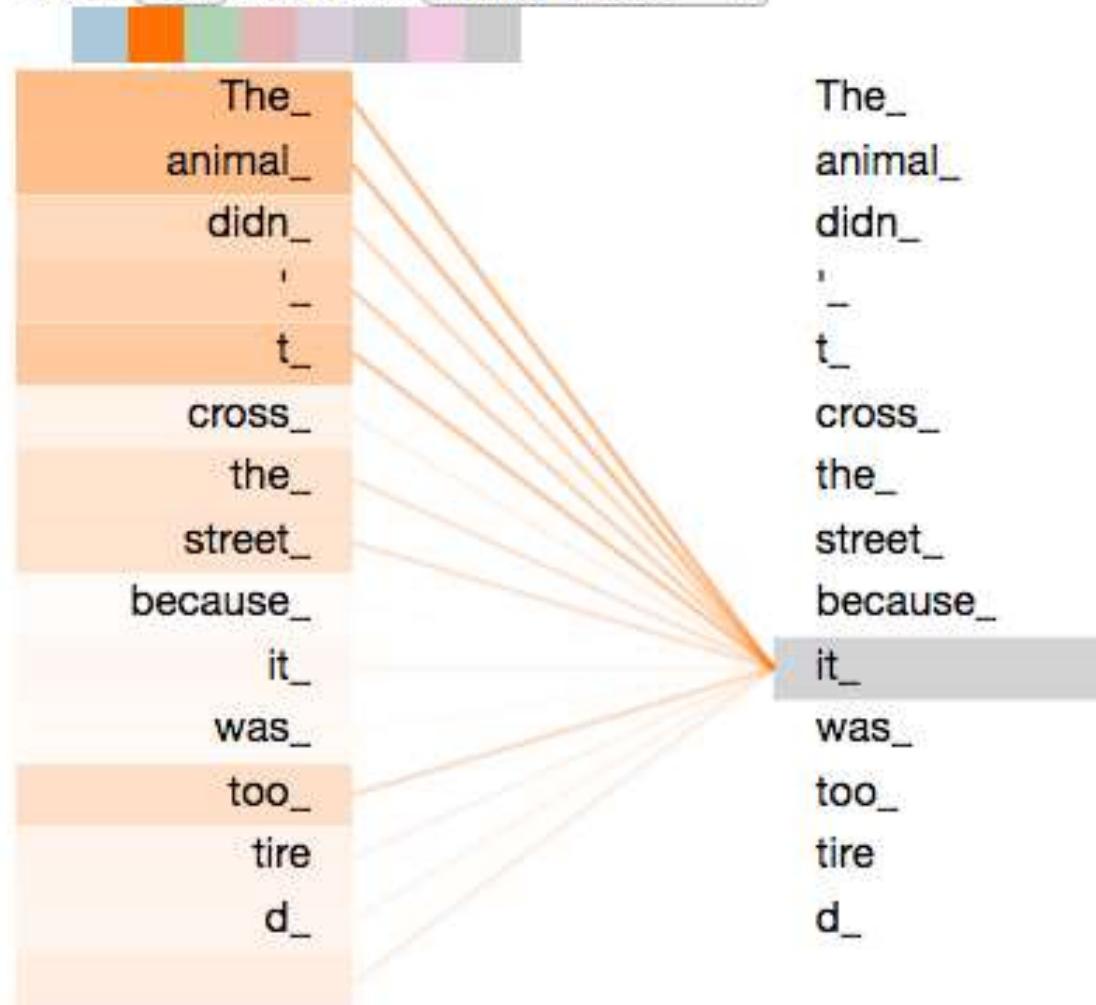
- When reading the sentence:

“The animal didn’t cross the street because it was too tired.”
To understand “it”, we must look at “animal”, not “street”.
- Self-attention lets each word pull in information from other relevant words.

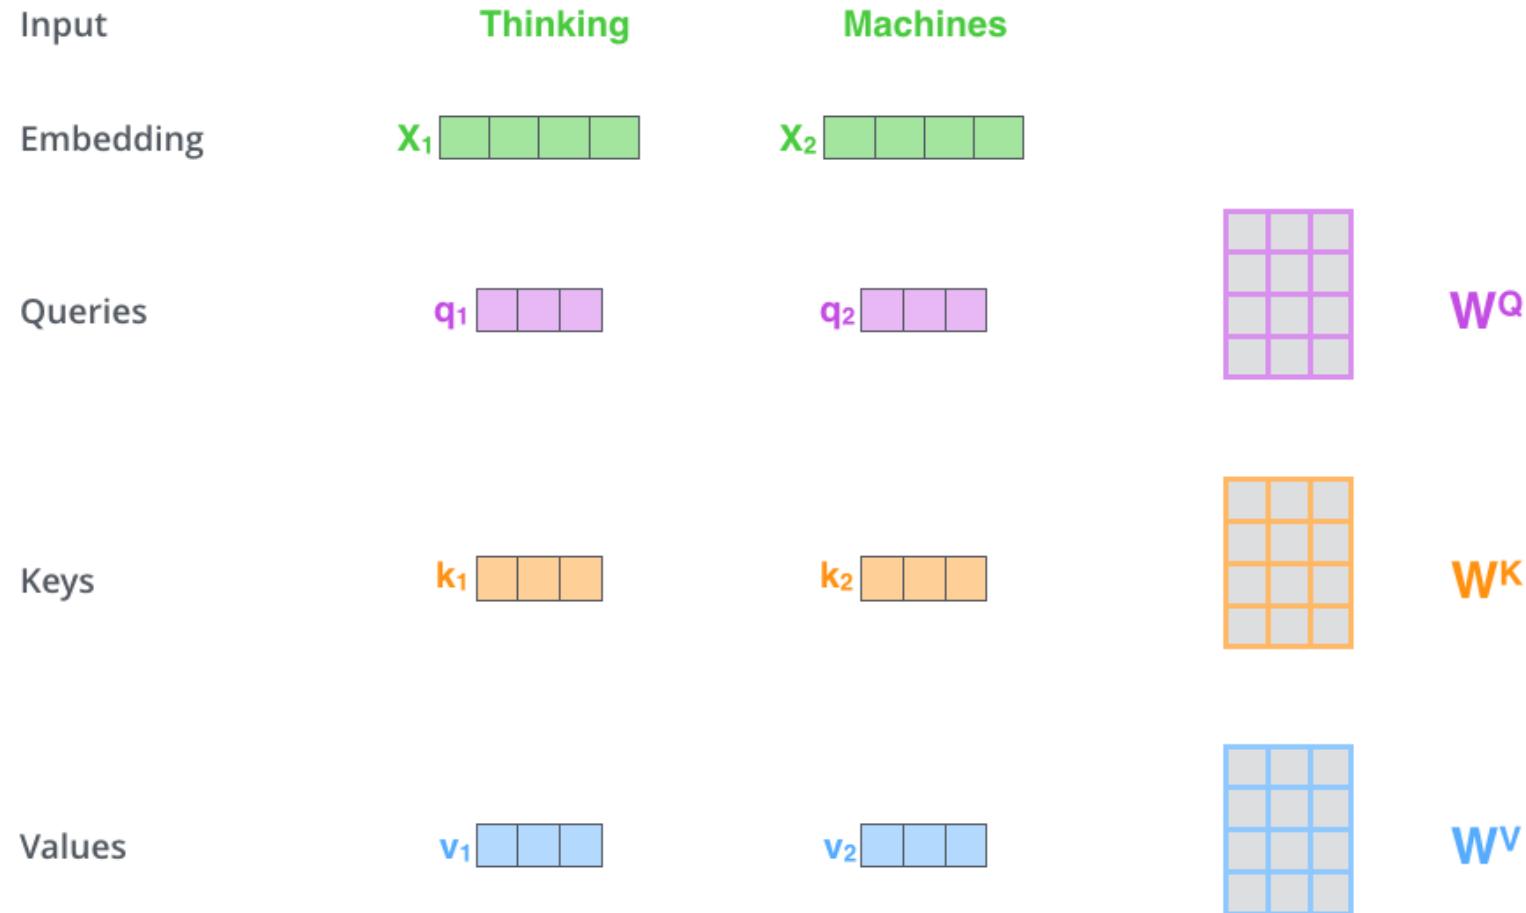
For each word, self-attention asks:

- “Which words in this sentence should I pay attention to?” and
- “How much should I pay attention to each one?”

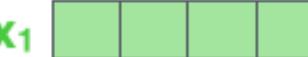
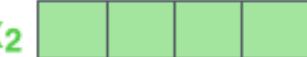
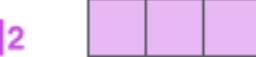
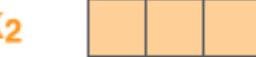
Layer: 5 Attention: Input - Input



Self-Attention in Detail



The first step in calculating self-attention is to create three vectors from each of the encoder's input vectors.

Input		
Embedding	Thinking	Machines
Queries	x_1 	x_2 
Keys	q_1 	q_2 
Values	k_1 	k_2 
Score	$q_1 \cdot k_1 = 112$	
	$q_1 \cdot k_2 = 96$	

The second step in calculating self-attention is to calculate a score.

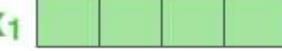
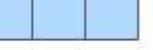
Interpretation of Score value

- High score → The second word is very relevant to the first
- Low score → The word is not important
- Near zero → The words are unrelated
- Negative (sometimes) → They should be ignored

So while computing attention for the query of “it”, we expect:

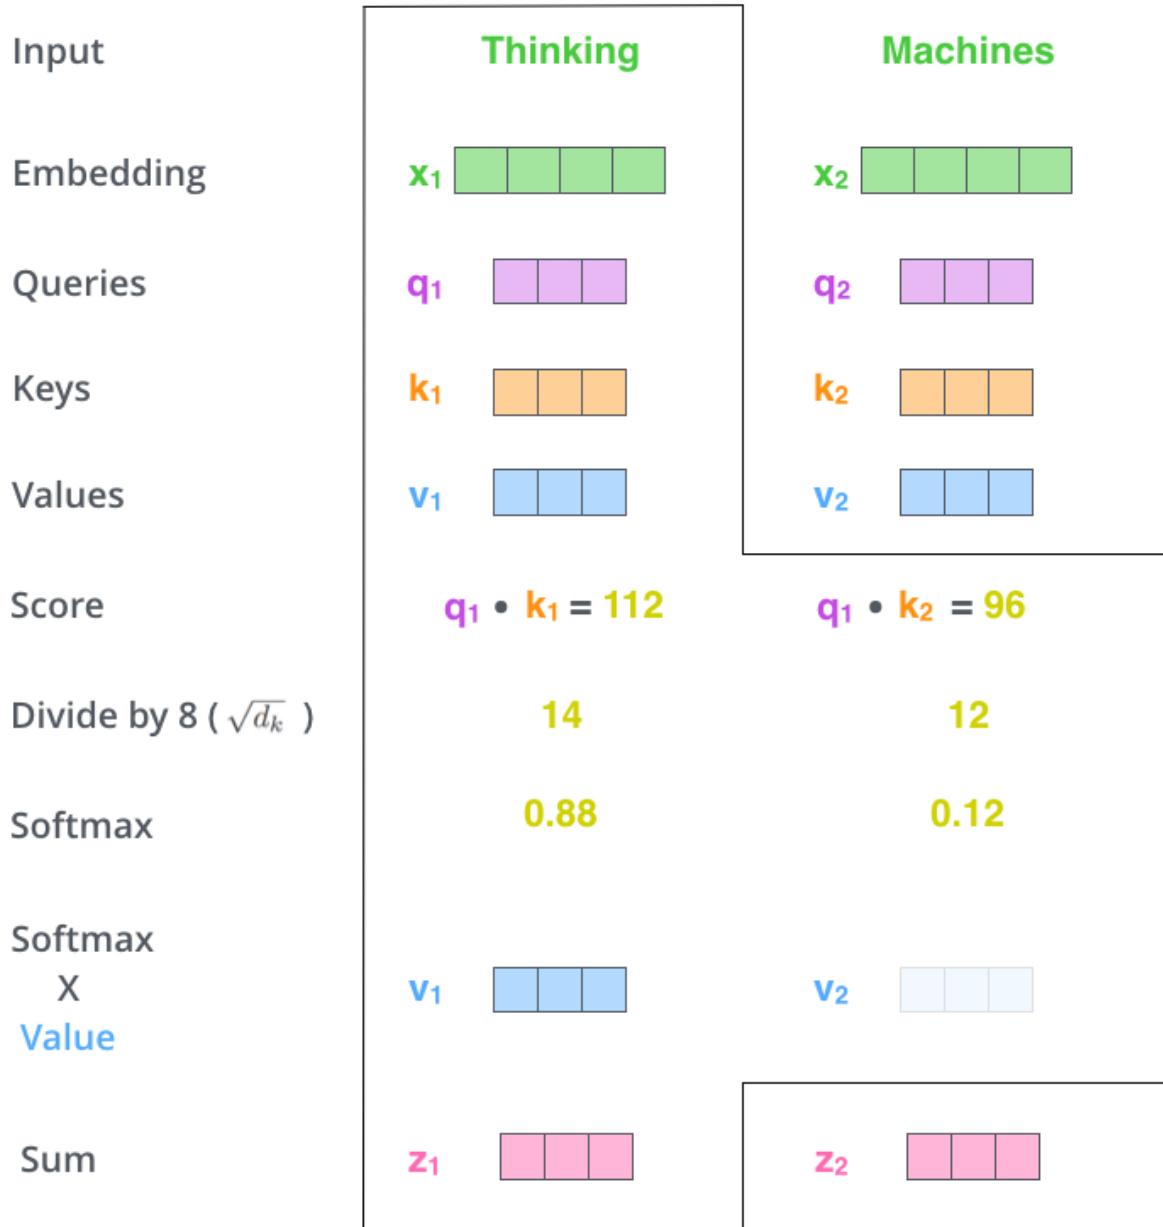
- $q_{it} \cdot k_{animal}$ → highest score
- $q_{it} \cdot k_{street}$ → lower
- $q_{it} \cdot k_{it}$ → moderate
- $q_{it} \cdot k_{otherwords}$ → lower

In one sentence: The attention score ($q \cdot k$) is a learned measure of how much one word should pay attention to another.

Input		
Embedding	x_1	
Queries	q_1	
Keys	k_1	
Values	v_1	
Score	$q_1 \cdot k_1 = 112$	$q_1 \cdot k_2 = 96$
Divide by 8 ($\sqrt{d_k}$)	14	12
Softmax	0.88	0.12

Step 3 is to divide the score by 8(square root of dimension of key vectors) and step 4 is to use softmax operation.

- Why divide by 8?
 - Without dividing, the dot product becomes too large and can cause exploding gradient problems(training becomes unstable)
- Why softmax?
 - Softmax takes raw scores and turns them into a probability distribution.
 - That means:
 - all values become positive
 - they add up to 1
 - bigger scores become bigger weights
 - smaller scores become smaller weights
 - It answers the question: “How much attention should I give to each word?”



Fifth step is to multiply each value vector by softmax score. And 6th step is to sum up the weighted value vectors.

- Why multiply?

Because the Value vectors represent the actual content of each word.

- High score → keep more of that word's content
- Low score → keep less of it
- Zero score → ignore that word completely

- And then finally add all the value vectors to give Z which is the output of self attention.
- The Z vector contains:
 - The word itself
 - Plus relevant information from other words
 - Weighted by how important they are

Matrix Calculation of Self-Attention

$$\mathbf{X} \times \mathbf{W}^Q = \mathbf{Q}$$

A diagram illustrating the calculation of the Query matrix (\mathbf{Q}) from the input matrix (\mathbf{X}). It shows a green 3x4 matrix \mathbf{X} multiplied by a purple 4x4 weight matrix \mathbf{W}^Q , resulting in a purple 3x4 matrix \mathbf{Q} . The matrices are represented as grids of colored squares.

$$\mathbf{X} \times \mathbf{W}^K = \mathbf{K}$$

A diagram illustrating the calculation of the Key matrix (\mathbf{K}) from the input matrix (\mathbf{X}). It shows a green 3x4 matrix \mathbf{X} multiplied by an orange 4x4 weight matrix \mathbf{W}^K , resulting in an orange 3x4 matrix \mathbf{K} . The matrices are represented as grids of colored squares.

$$\mathbf{X} \times \mathbf{W}^V = \mathbf{V}$$

A diagram illustrating the calculation of the Value matrix (\mathbf{V}) from the input matrix (\mathbf{X}). It shows a green 3x4 matrix \mathbf{X} multiplied by a blue 4x4 weight matrix \mathbf{W}^V , resulting in a blue 3x4 matrix \mathbf{V} . The matrices are represented as grids of colored squares.

$$\text{softmax}\left(\frac{\mathbf{Q} \times \mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V}$$

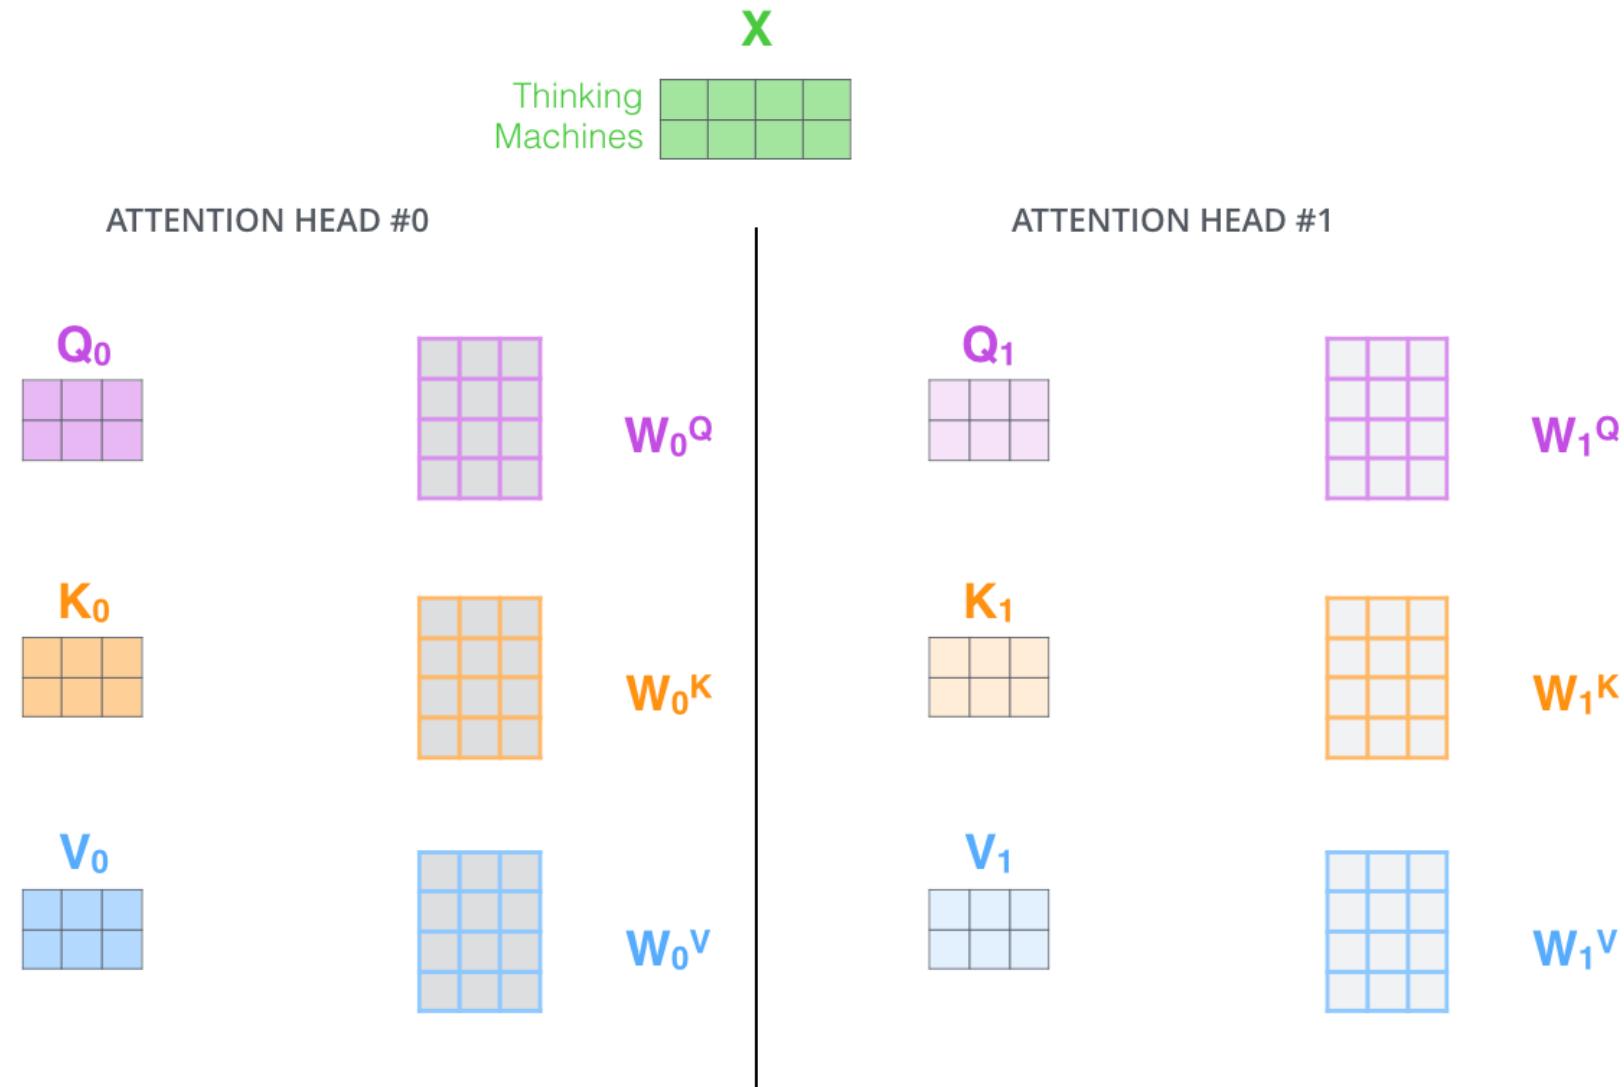
= \mathbf{Z}

The diagram illustrates the self-attention calculation in matrix form. It shows the softmax function taking the product of \mathbf{Q} and \mathbf{K}^T (scaled by $\sqrt{d_k}$) as input, and multiplying the result by \mathbf{V} . Below, an equals sign is followed by a pink 2x3 matrix labeled \mathbf{Z} .

The self-attention calculation in matrix form.

Multi-Head Attention

- Why multi-head attention?
 - One head cannot capture all the patterns at once.
 - Multiple heads allow the transformer to understand language from many angles at the same time.
- In the example sentence:
 - “The animal didn’t cross the street because it was tired.”, different kinds of relationships matter.
 - Head 1: “it” \leftrightarrow “animal” (pronoun resolution)
 - Head 2: grammatical structure
 - Head 3: adjective \rightarrow noun and so on.



X

Thinking
Machines

Calculating attention separately in
eight different attention heads

ATTENTION
HEAD #0

z_0

ATTENTION
HEAD #1

z_1

...

ATTENTION
HEAD #7

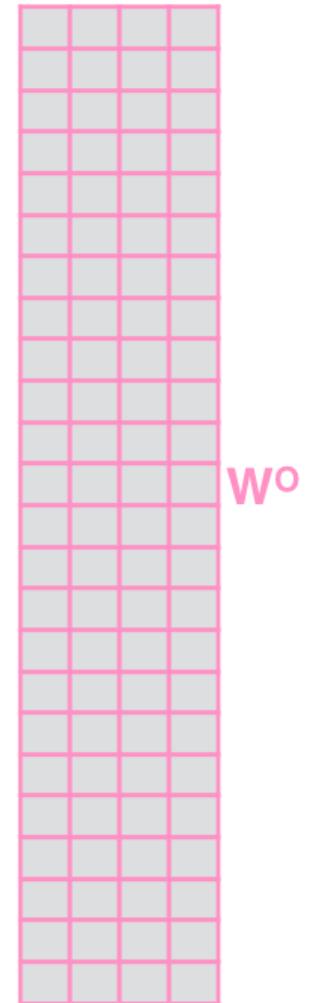
z_7

1) Concatenate all the attention heads



2) Multiply with a weight matrix W^o that was trained jointly with the model

X

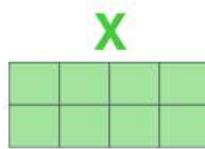


3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN

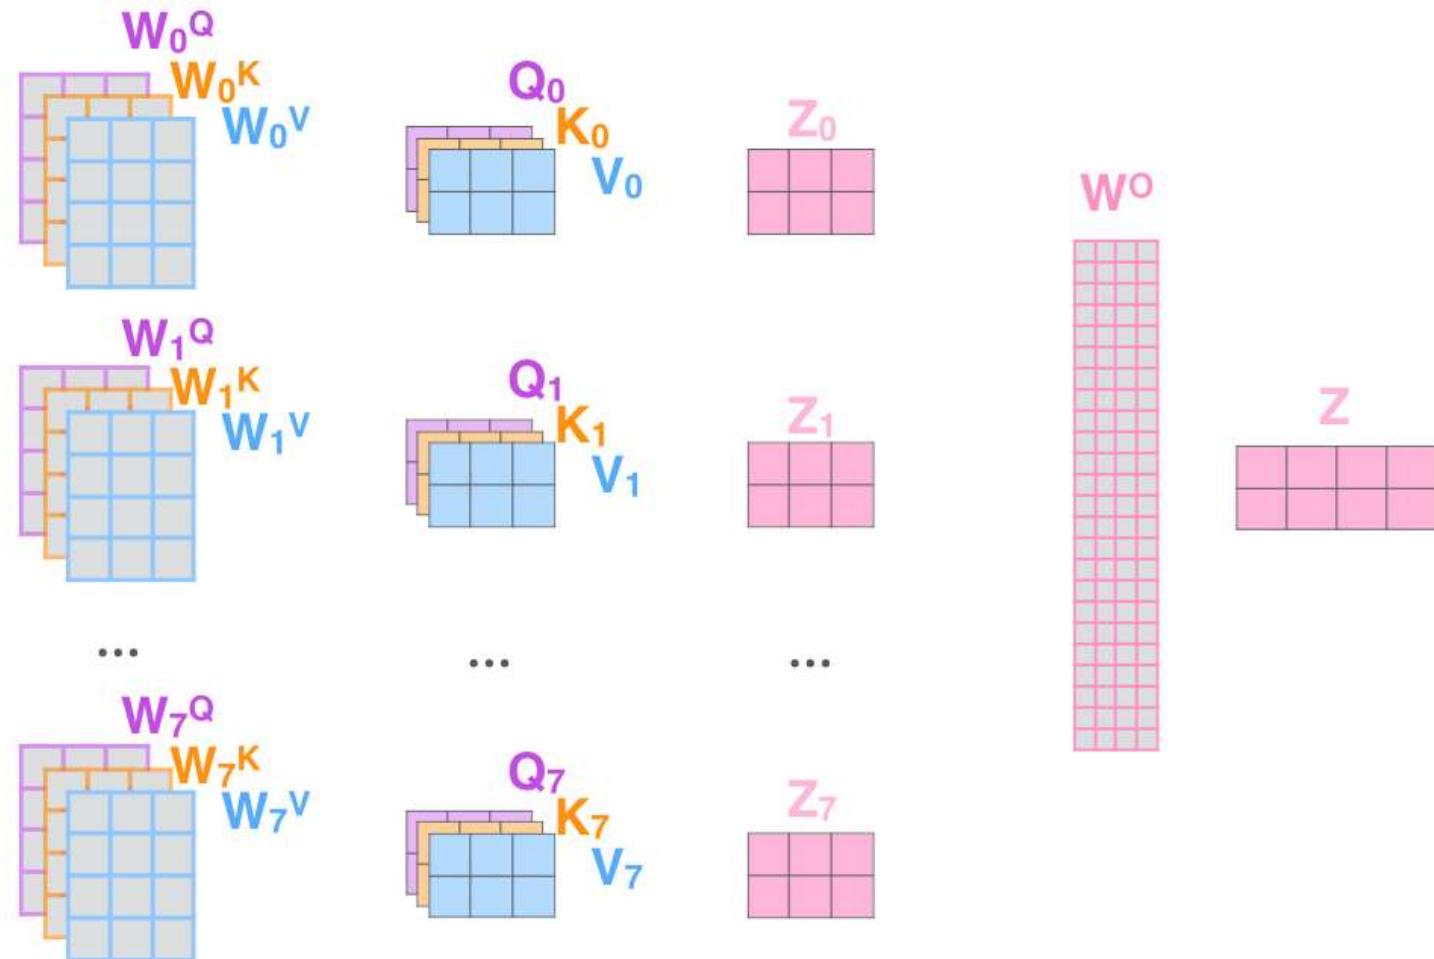
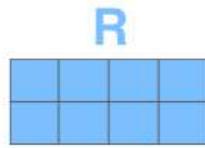
$$= \begin{matrix} Z \\ \hline \text{---} \\ \text{512d} \end{matrix}$$

- 1) This is our input sentence* X
- 2) We embed each word* R
- 3) Split into 8 heads. We multiply X or R with weight matrices
- 4) Calculate attention using the resulting $Q/K/V$ matrices
- 5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

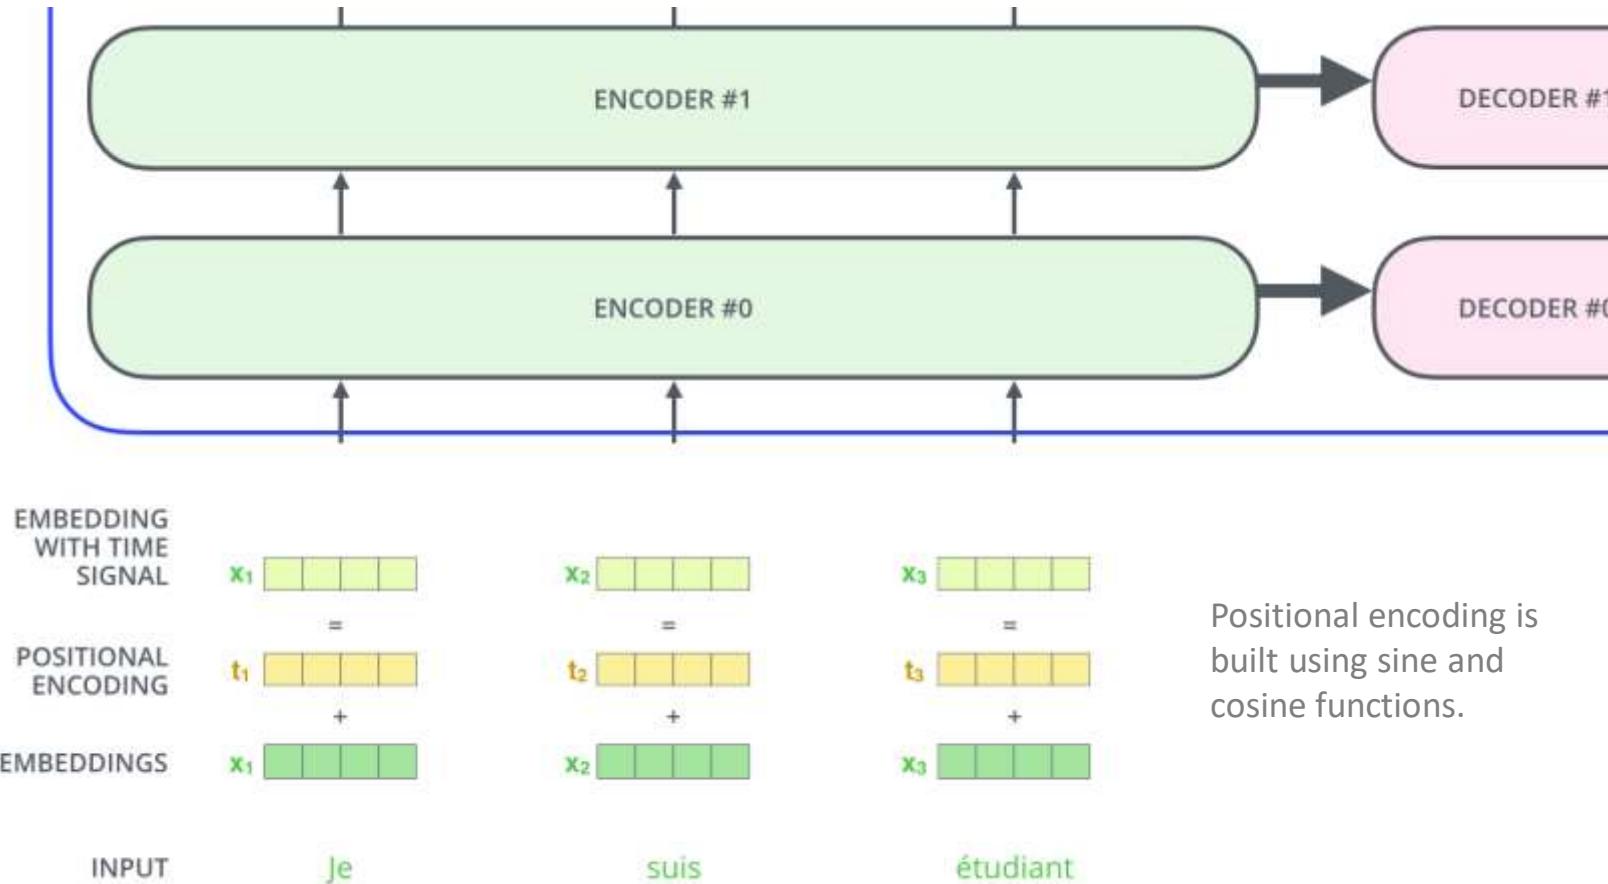
Thinking
Machines



* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

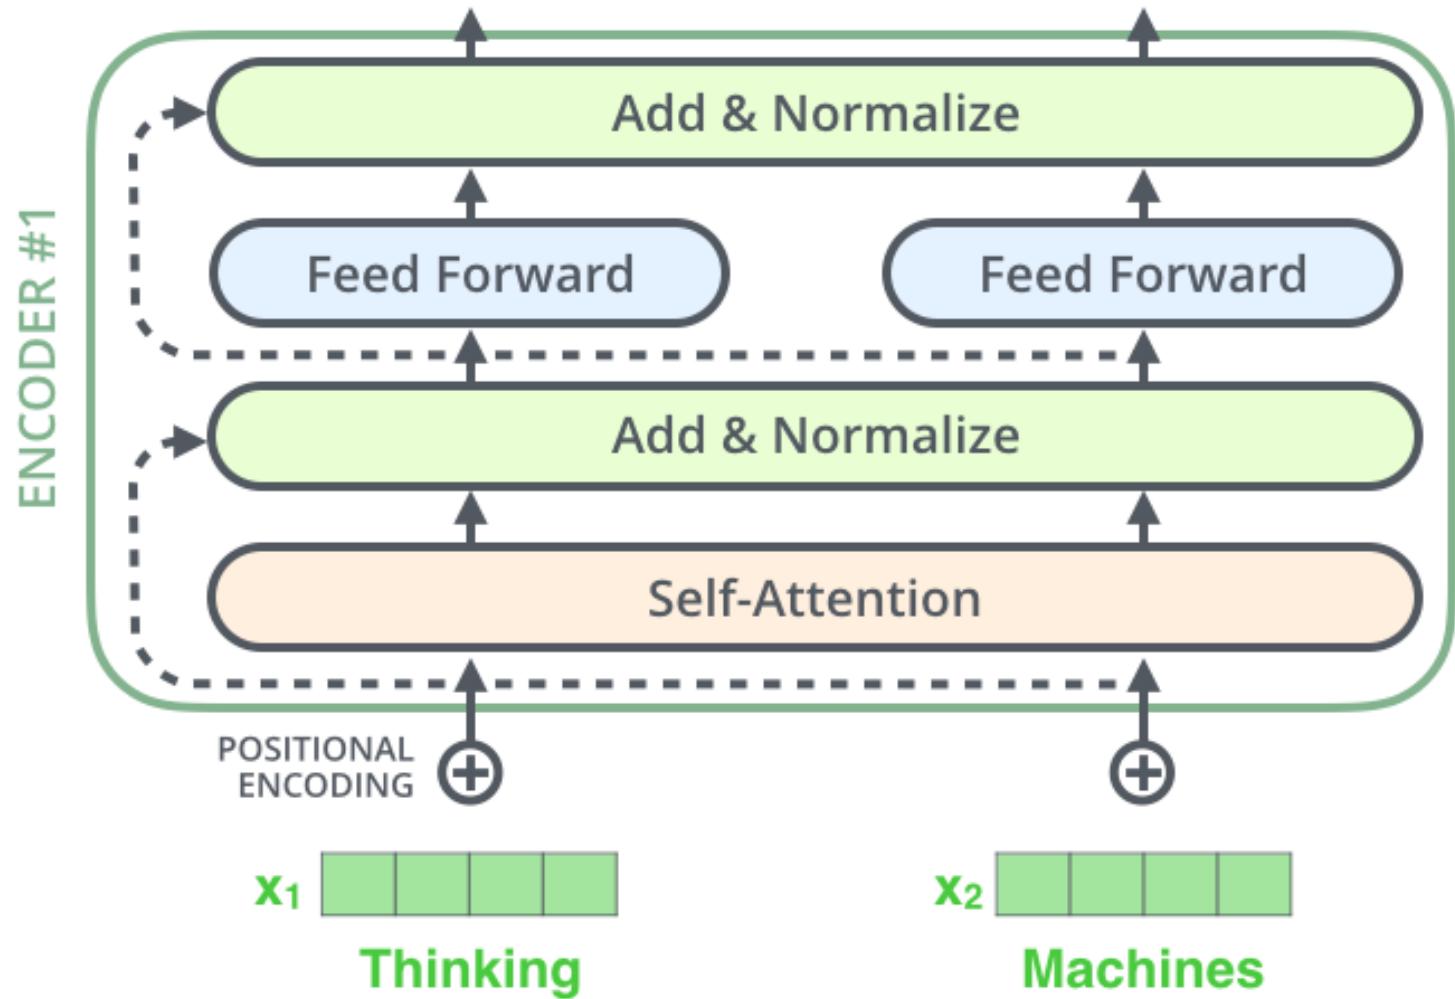


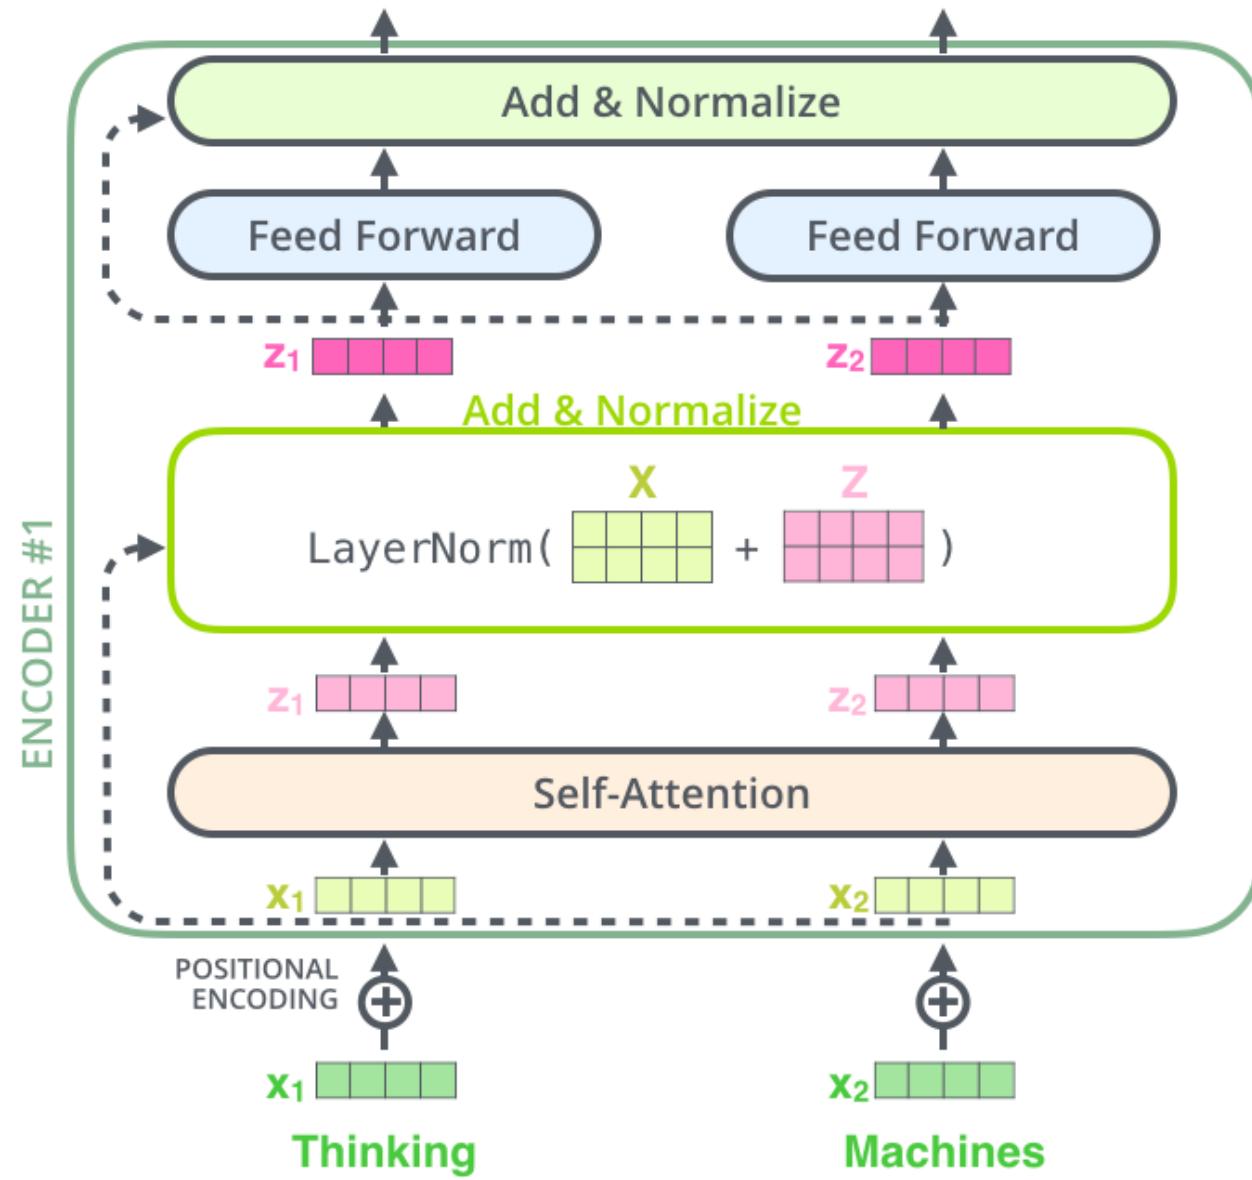
Representing The Order of The Sequence Using Positional Encoding

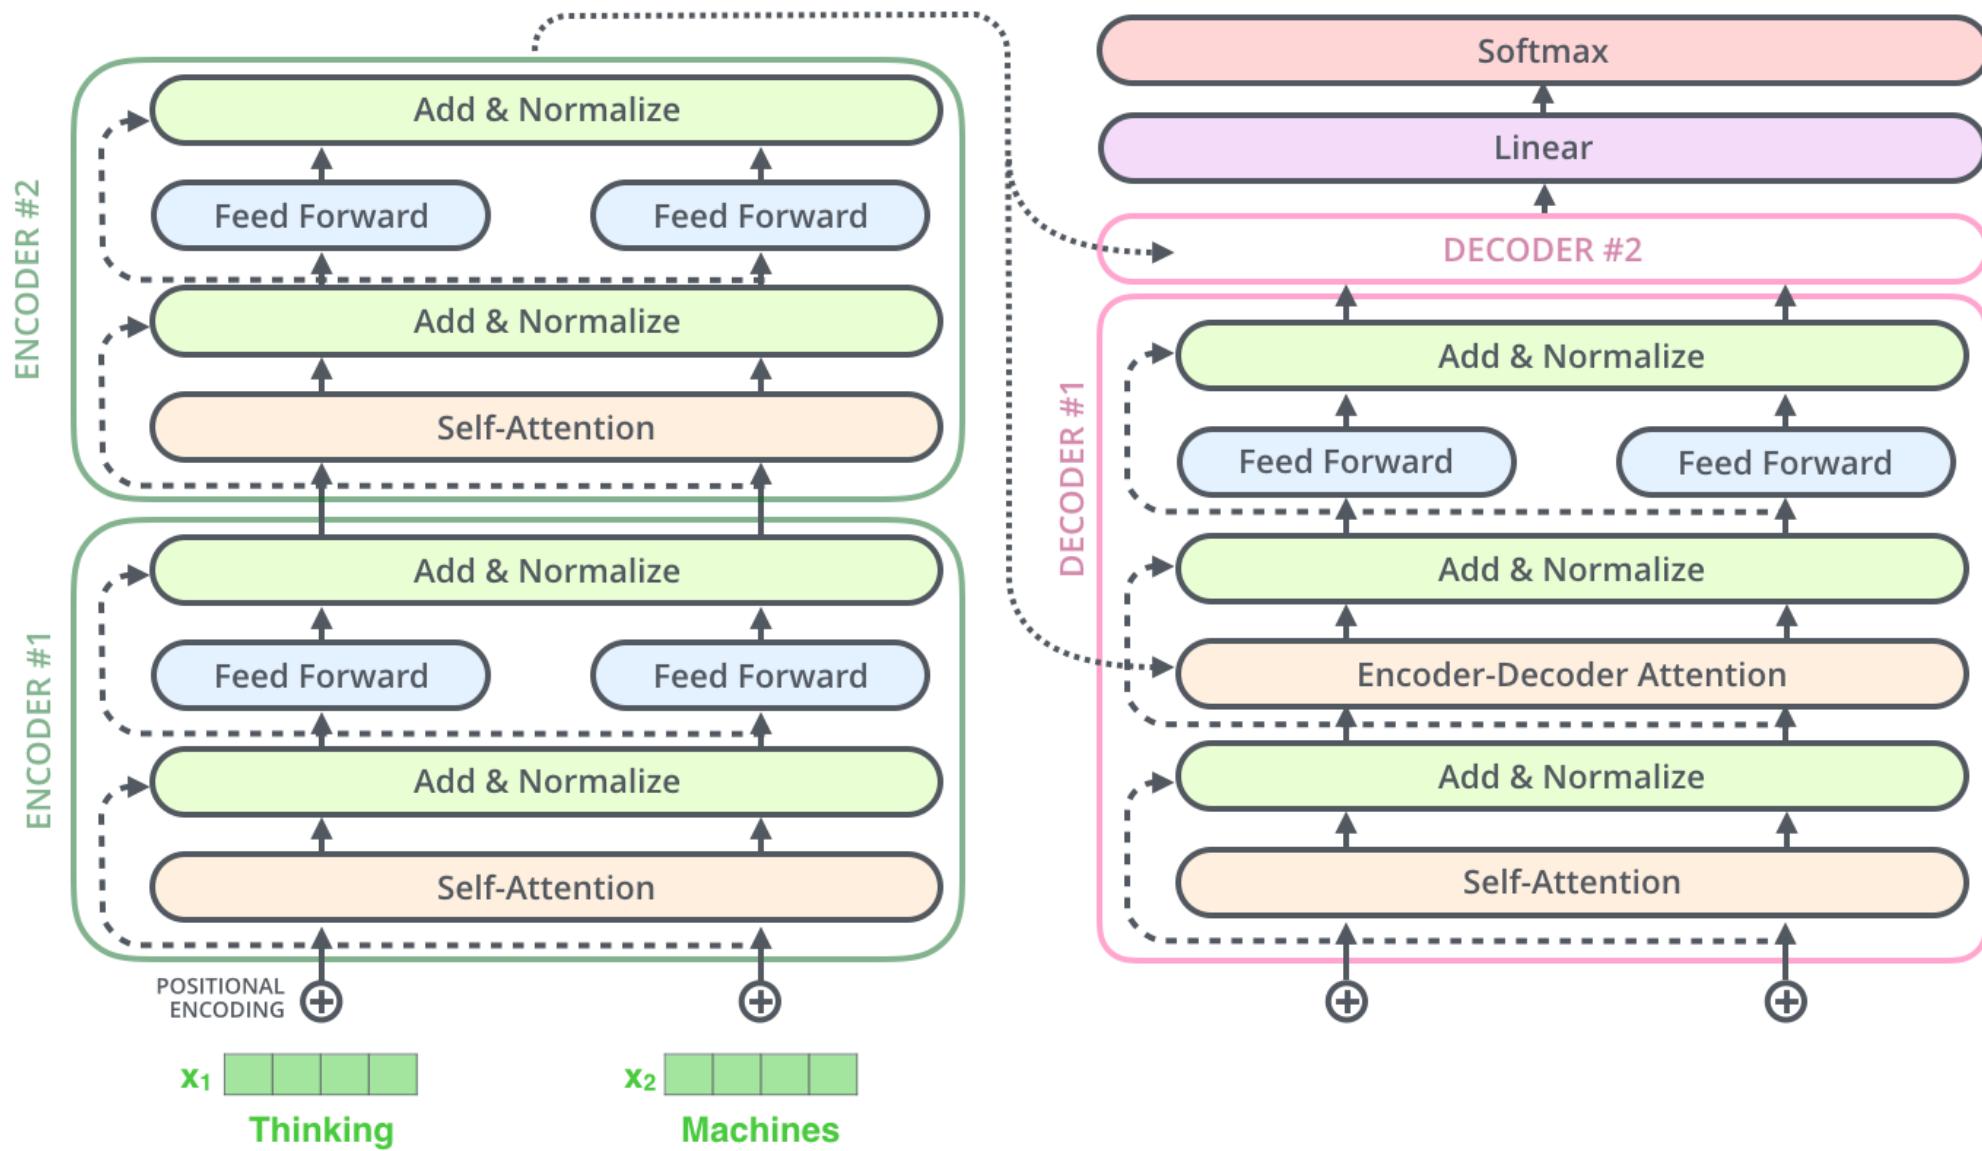


The Residuals

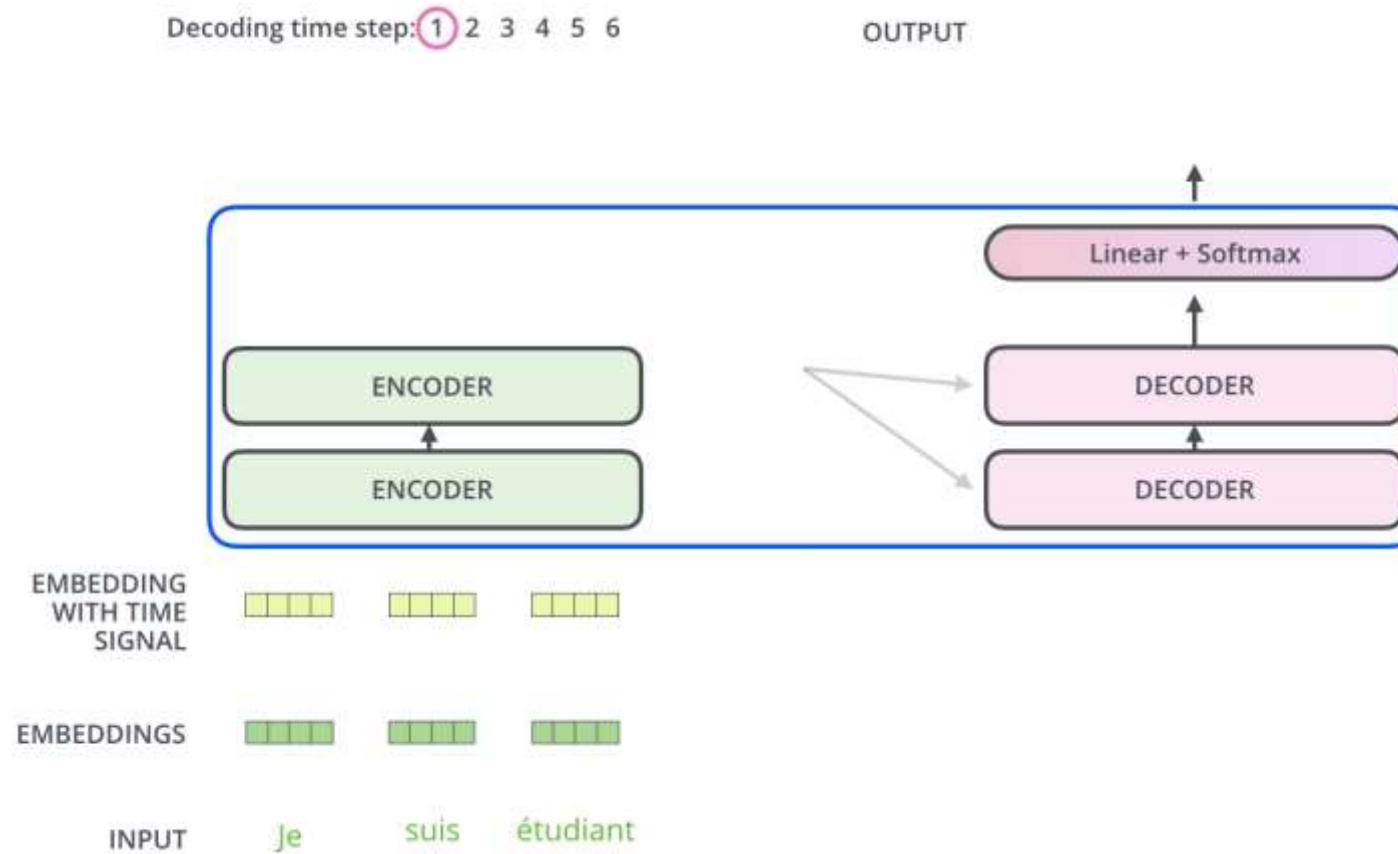
- It is a skip connection where the original input is added back to the output of:
 - Self-attention
 - Feed-forward network
- Why do we need residuals?
 - To avoid vanishing gradient problems.





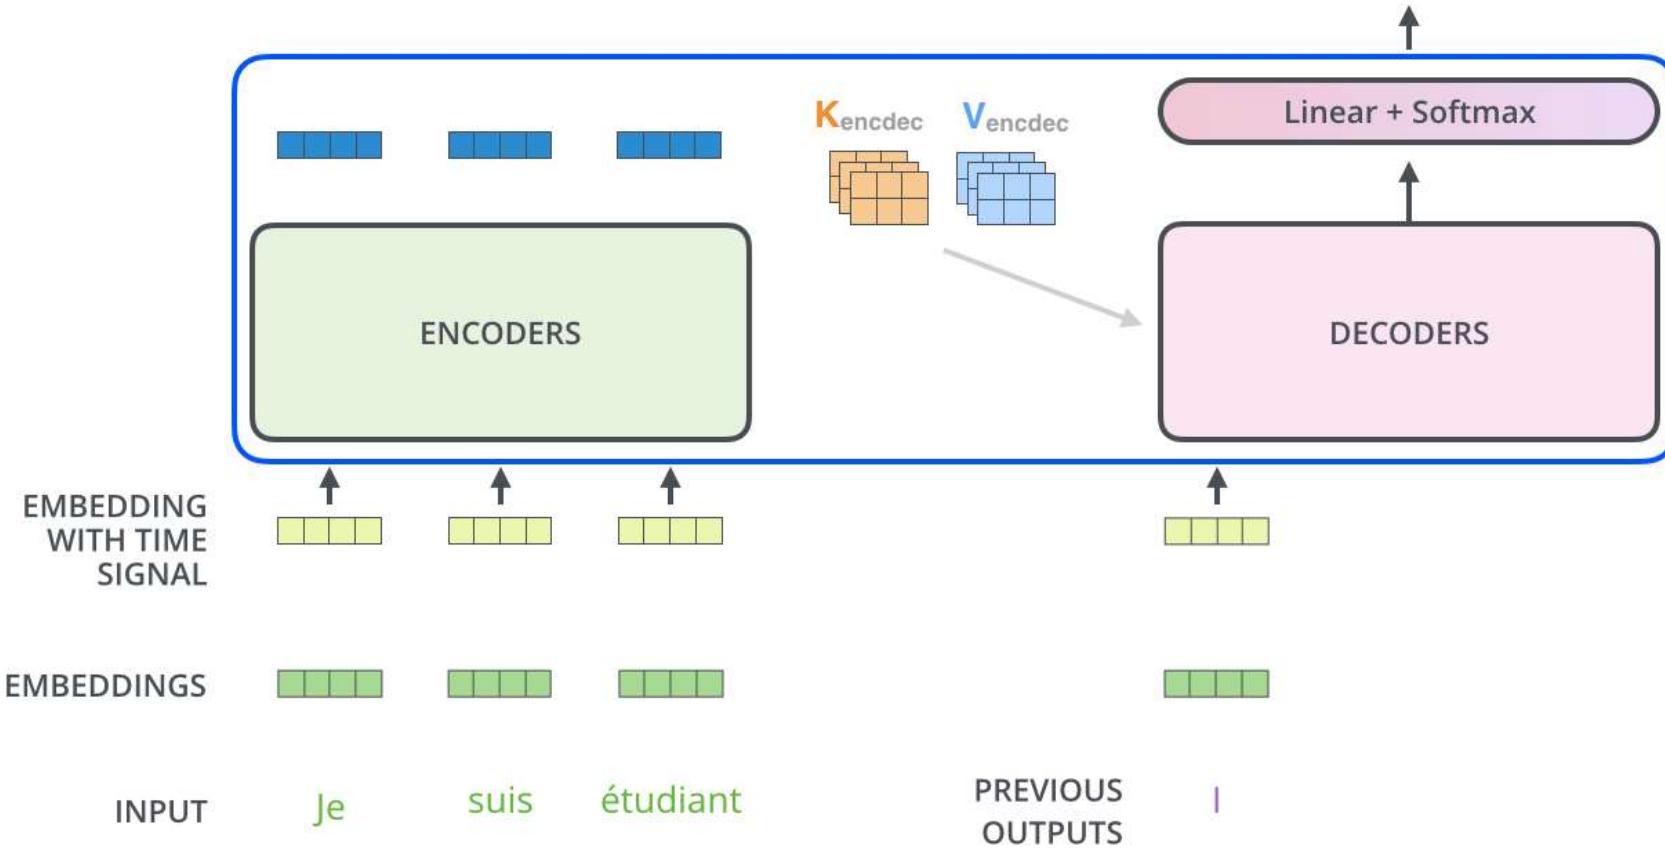


The Decoder Side

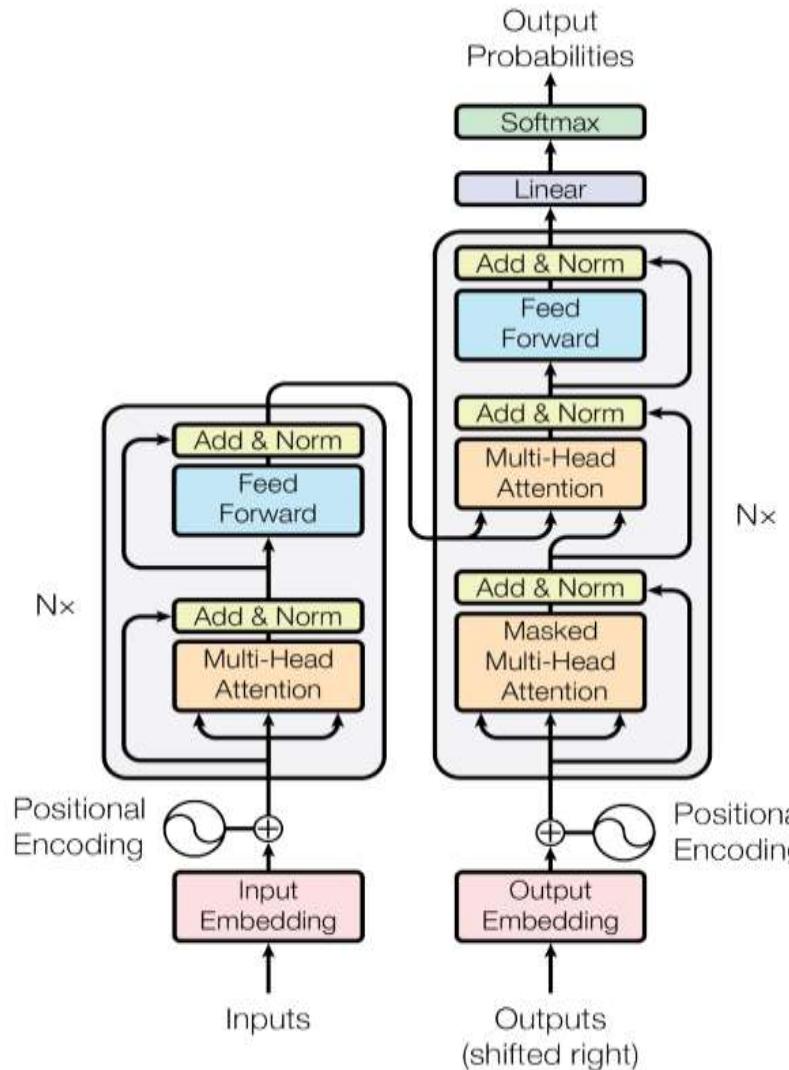


Decoding time step: 1 2 3 4 5 6

OUTPUT |



The architecture from the paper



Reference

- <https://jalammar.github.io/illustrated-transformer/>

Thank You!